



DEPARTMENT OF ELECTRIC ENERGY

TET4510 - SPECIALIZATION PROJECT

Short-Term mFRR Activation Direction Forecasting at 15-Minute Resolution

Author:
Haakon Nygård Hellebust

Date

Table of Contents

List of Figures	ii
List of Tables	iii
1 Introduction	1
1.1 Background	1
1.2 Motivation	1
1.3 Research Question	2
1.4 Outline	2
2 Theory	2
2.1 Electricity Balancing Market	2
2.1.1 Nordic Balancing Market Structure	3
2.1.2 Reserve Market Concepts	4
2.2 Machine Learning Theory	5
2.2.1 Supervised Learning for Time-Dependent Classification	5
2.2.2 Class Imbalance	5
2.2.3 Tree-Based Models	5
3 Literature Review	6
3.1 LOOKING AT PROPOSED ACTIVATION UNCERTAINTY VARIANTS	6
3.2 mFRR Energy Activation Market Characteristics	6
3.3 Existing Approaches to Balancing Activation Market Forecasting	7
3.3.1 Imbalance Volume and Price Forecasting	7
3.3.2 Activation Forecasting	8
3.3.3 Literature Synthesis	8
3.3.4 Research Gap	9
4 Methodology	10
4.1 Overview of Methodological Approach	10
4.2 Data Sources	10
4.2.1 Nord Pool	10
4.2.2 NUCS	11
4.2.3 ENTSO-E	11
4.3 Data Preprocessing and Analysis	12
4.3.1 Dataset structure	12

4.3.2	Resampling, Imputation, and Merging	12
4.3.3	Time Restrictions	13
4.3.4	Exploratory Data Analysis	14
4.4	Feature Engineering	17
4.4.1	Lag features	17
4.4.2	Cross-zonal flow features	18
4.4.3	Temporal features	19
4.4.4	Price features	19
4.4.5	Load features	19
4.4.6	Production features	19
4.4.7	Interaction features	20
4.5	Evaluation Framework	20
4.5.1	Data Splitting	23
4.5.2	Model Selection	23
4.5.3	Adjusting Classification Thresholds	23
4.5.4	Decision-bias tuning	24
4.5.5	AutoGluon	25
5	Model Results	25
5.1	Naive Model	25
5.2	The 2025 Dataset	25
5.2.1	Extra Trees	25
5.3	The 2024-2025 Dataset	27
5.3.1	CatBoost Models	28
5.4	Correlation	30
Bibliography		31

List of Figures

1	Norway electricity supply by power station and net imports with projections to 2050 [5]	1
2	The Nordic balancing market hierarchy, illustrating the different reserve types and their activation times.	3
3	Illustration of the different reserve types and their activation times.	4
4	Overview of the methodological approach used in this project.	10
5	Visualization of resampling from 1-hour to 15-minute resolution using forward filling. .	12

6	mFRR activation distribution.	14
7	Cross-zonal flow distributions for the NO1 bidding zone.	15
8	Cross-zonal flow utilizations for the NO1 bidding zone calculated as the ratio between actual flow and NTC capacity.	15
9	A histogram of hourly aFRR procurement prices for the NO1 bidding zone from NUCS data.	16
10	Hourly aFRR procurement prices for the NO1 bidding zone from NUCS data.	16
11	Load and production data distributions.	20
12	Example of a confusion matrix for a three-class classification problem.	21
13	Typical Precision-Recall Curve with F1-score maximization point indicated.	24
14	Precision-Recall Curve for Extra Trees model trained on 2025 dataset with highest F1-score.	26
15	Feature importance for Extra Trees model trained on 2025 dataset with highest F1-score.	27
16	Feature importance for CatBoost model trained on 2024-2025 dataset after hyper-parameter tuning.	28
17	Performance metrics for CatBoost model on 2024-2025 dataset at different thresholds.	29
18	Up/down price minus day-ahead price distribution for CatBoost model trained on 2024-2025 dataset. quick_multiclass_cat_hpo	29
19	Performance metrics for CatBoost model on 2025 March 4th dataset at different confidence thresholds.	30

List of Tables

1	Comparison of key studies on balancing market forecasting	9
2	Summary statistics for wind-related features (2024–2025, NO1)	16
3	Performance metrics for Extra Trees model on 2025 dataset at different thresholds.	26

1 Introduction

1.1 Background

The increasing share of weather-dependent renewable energy sources has reduced the inertia of modern power systems and made short-term system imbalances more frequent and difficult to predict [1]. To maintain frequency quality, Nordic transmission system operators (TSO) rely on various balancing resources, among which manual Frequency Restoration Reserves (mFRR) play a key role. The mFRR energy activation market (mFRR EAM) enables market participants to offer up- and down-regulation reserves that the TSO can activate when sustained imbalances occur [2].

Historically, mFRR activations were manual and scheduled on an hourly basis. On 4 March 2025, the Nordic region introduced a new automatic mFRR activation platform with a 15-minute resolution [3]. This reform aligns the Nordic system with European balancing-market standards and allows activation signals to more closely follow real-time system conditions. Automated activation of mFRR reserves enables faster response times, making it possible to address quarter-hourly imbalances effectively [4].

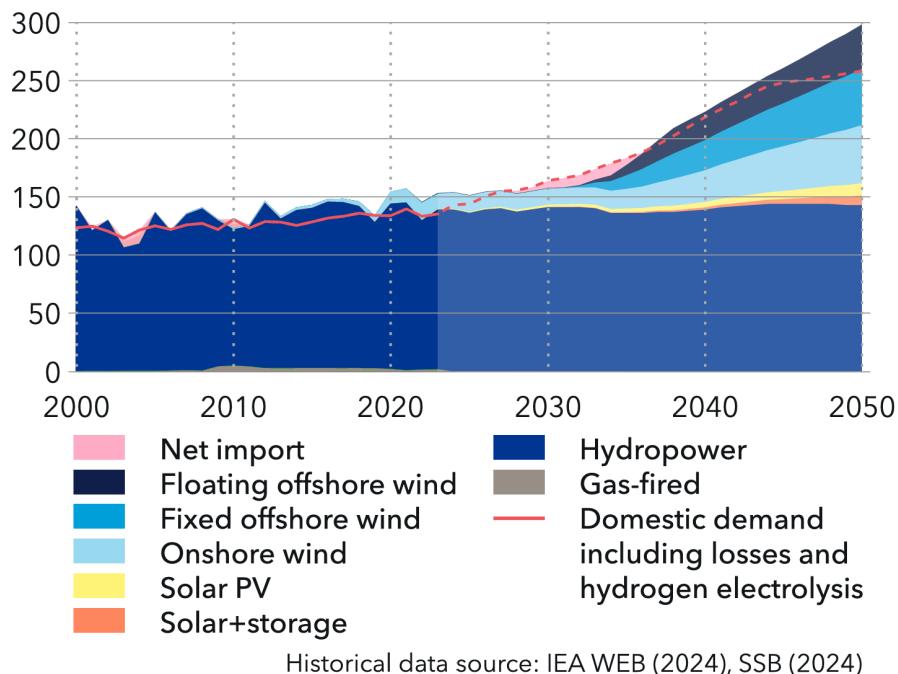


Figure 1: Norway electricity supply by power station and net imports with projections to 2050 [5]

Struggling with what to include in background as I feel like much of interesting background information kind of fits in motivation.

1.2 Motivation

mFRR activations have direct economic implications for market participants. Accurately anticipating whether up- or down-regulation is likely enables participants to choose between committing flexible resources to capacity markets or bidding directly into the activation market, where revenues depend on actual activations. Capacity markets provide stable income, whereas activation markets offer potentially higher but uncertain returns. Being able to accurately assess the likelihood of an activation therefore supports more informed and profitable bidding strategies.

For aggregators of distributed flexible resources, such as electric vehicles, heat pumps, and batteries,

predicting mFRR activations is particularly valuable. When committing flexibility to the mFRR capacity market, the aggregator earns a fixed payment for being available to provide reserves. However, this also means the resources are locked up and cannot be used for other purposes during the committed period. Participation in the activation market locks resources for shorter periods, allowing the aggregator to remain flexible. The guaranteed income from capacity payments must thus be weighed against the opportunity cost of missing out on potential revenues from the activation market. The correct choice depends on the likelihood of activations occurring, which is inherently uncertain.

A reliable activation prediction model can substantially reduce this uncertainty by providing an predictions of whether up- or down-regulation is likely in future intervals. With such information, an aggregator can manage its portfolio more strategically, for example, by reserving flexibility for the activation market only when the probability of activation is sufficiently high, or by safely committing to the capacity market when activations appear unlikely. Accurate predictions therefore enable more efficient utilization of distributed resources and support more economically rational market participation.

The recent shift to 15-minute resolution further increases the relevance of such models. While higher time resolutions offer greater responsiveness, it also creates a shorter decision window and increases sensitivity to rapidly changing system conditions. This motivates data-driven prediction methods as an approach to support market bidding strategies. An interesting application is to use the predictions to inform reinforcement learning-based bidding strategies for aggregators of distributed flexible resources. Reinforcement learning agents can learn to make optimal decisions, but require a reliable model of the environment to base their strategies on. A well-performing mFRR activation prediction model could serve as this model of the environment, enabling the development of near-optimal bidding strategies that adapt to predicted market conditions. **Is this too much reinforcement learning talk for an introduction?**

1.3 Research Question

The central research question of this study is: How well can mFRR up- and down-regulation activations be predicted at 15-minute resolution using only real-time-available system state features and historical activation/price data? **Change this I suppose?**

1.4 Outline

This report is structured as follows: Chapter 2 provides an overview of the Nordic balancing markets and related work. Chapter 3 reviews existing literature on imbalance and activation forecasting and identifies the research gaps addressed in this study. Chapter 4 details the modelling methodology, including data preprocessing, feature engineering, and model evaluation and selection. Chapter 5 presents the results, while Chapter 6 discusses implications, limitations, and directions for future work.

2 Theory

2.1 Electricity Balancing Market

Electricity balancing markets are mechanisms designed to ensure the stability and reliability of the power grid by managing supply and demand in real-time. These markets facilitate the procurement of balancing services, which are necessary to maintain the equilibrium between electricity supply and demand. Balancing markets operate on the principle of economic efficiency, where market participants can offer their flexibility to the grid operator in exchange for compensation.

2.1.1 Nordic Balancing Market Structure

In Nordic countries, the balancing hierarchy consists of several layers, each serving a specific purpose in maintaining grid stability. Figure 2 illustrates the different reserve types, their activation times, and their place in the hierarchy.

Not manual anymore, I will make my own figure.

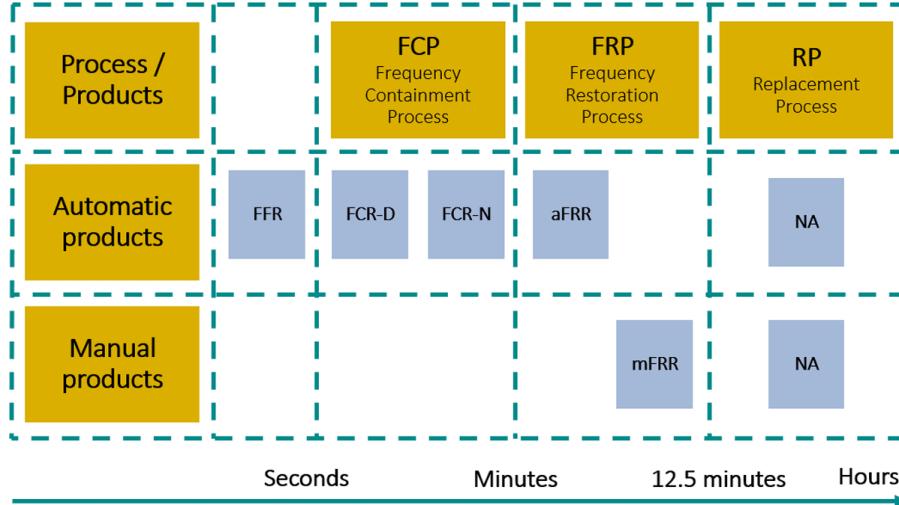


Figure 2: The Nordic balancing market hierarchy, illustrating the different reserve types and their activation times [2].

Frequency Containment Reserves (FCR) were designed to be the first line of defense against frequency deviations in the power grid. These reserves are activated automatically and respond quickly to counteract sudden imbalances between supply and demand. FCR is further divided into two categories: FCR-N and FCR-D. FCR-D is specifically intended to address frequency deviations caused by disturbances in the distribution network, while FCR-N focuses on normal operating conditions in the transmission network. FCR-D should, therefore, be able to respond faster than FCR-N to effectively manage these disturbances.

The Fast Frequency Reserves (FFR) reserve market was implemented in the Nordics in May 2020. These reserves are designed to respond even more rapidly than FCR, ideally in the span of a single second. The need for FFR arises from the increasing penetration of renewable energy sources. Wind power is, for instance, not connected synchronously to the grid, leading to a reduction in system inertia. Lower inertia means that frequency deviations occur more rapidly, necessitating faster-acting reserves like FFR to maintain grid stability. The fast power response provided by FFR is usually sustained for a short duration, stabilizing the frequency slightly before FCR-D takes over [6]. Figure 3 illustrates roughly the activation times and the interplay between the different reserve types in the Nordic balancing market.

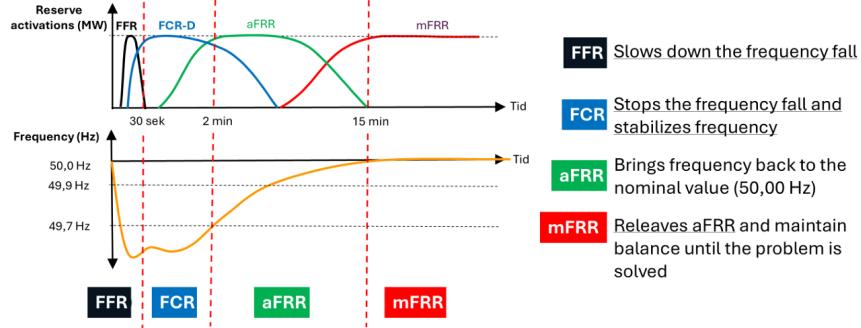


Figure 3: Illustration of the different reserve types and their activation times [2].

After FFR and FCR has stabilized the frequency, something must bring it back to its nominal level. Automatic Frequency Restoration Reserves (aFRR) holds this responsibility. These reserves are often kept activated for a couple of minutes to ensure that the frequency is restored to its normal operating level. Manual Frequency Restoration Reserves (mFRR) then relieves aFRR and maintains the balance until normal operations are restored.

2.1.2 Reserve Market Concepts

Reserve markets are platforms where market participants can offer their balancing services to the grid operator. These markets operate on the principle of supply and demand, where participants can bid to provide reserves at specific prices. The markets are then cleared based on the bids received, ensuring that the most cost-effective resources are utilized to maintain grid stability. Reserve markets can be broadly categorized into two types: activation markets and capacity markets. Only aFRR and mFRR markets will be discussed further, as they are the most relevant for this study.

Capacity Markets Capacity is a market mechanism that ensures the availability of sufficient resources to maintain grid stability and reliability. Capacities are procured prior to real-time operation to guarantee the availability of balancing resources at the time of operation [2]. BSPs can offer their capacities to the grid operator through the Nordic aFRR capacity market or in the national (Statnett in Norway) mFRR capacity markets. Capacity market participants are compensated for making their resources available to the grid operator, regardless of whether their resources are activated or not. They are, however, obligated to deliver the offered capacity when called upon by the grid operator.

Energy Activation Markets (EAMs) Activation markets operate closer to real-time and are designed to procure balancing energy to address immediate imbalances in the power grid. In the Nordic region, the aFRR and mFRR activation markets serve this purpose. In the mFRR EAM in Norway, BSPs submit bids to Statnett at least 45 minutes before the activation period. These bids specify the amount of up- or down-regulation capacity the BSP is willing to provide and the corresponding price. Statnett then forwards the bids to the clearing algorithm Nordic Libra AOF, which provides the activation volumes for the operational quarter-hour. Statnett then activates the selected bids based on the activation volumes provided by Nordic Libra AOF [7].

The mFRR EAM underwent a significant reform March 4th 2025, transitioning from national mFRR activation systems to a Nordic-wide mFRR activation market. The previously hourly resolution was also changed to a quarter-hourly resolution. This resolution change allows for more precise balancing, introducing more nuanced price signals.

Imbalance Prices

mFRR Market in Norway (will refactor this or remove this) In Norway, the mFRR market plays a crucial role in maintaining the stability of the power grid. Market participants can offer their flexibility to the grid operator through both activation and capacity markets.

2.2 Machine Learning Theory

Machine learning (ML) is a subset of artificial intelligence (AI) that focuses on developing algorithms and statistical models that enable computers to perform specific tasks without explicit instructions. Instead, these algorithms learn from data, identifying patterns and making decisions based on the information provided. In this project, machine learning techniques are employed to predict the occurrence of mFRR activations in NO1 based on historical data and relevant system state indicators.

2.2.1 Supervised Learning for Time-Dependent Classification

Supervised learning is a machine learning paradigm where models are trained to map input data to specific outputs based on example input-output pairs. In this project, the input data consists of information about the current state relevant to the occurrence mFRR activations. The output is a ternary label indicating whether an up-regulation, down-regulation, or no activation occurs. Thus, the output variable is categorical, making this a classification task, where the goal is to assign input data points to one of several predefined classes. However, the problem at hand is not a standard supervised learning task, as the data is ordered and time-dependent. This is often referred to as *event prediction*, where *events* are defined as nontrivial occurrences in specific locations and time [8]. In this project, location is fixed to the NO1 bidding zone, while time is discretized into quarter-hourly intervals.

The chronological structure has wideranging implications for model training and evaluation. Standard supervised learning techniques often assume that data points are independent and identically distributed (i.i.d.), which is not the case for time-dependent data. Temporal dependencies must be accounted for, as past events can influence future outcomes. This calls for specialized model and data engineering techniques that can effectively capture and leverage these relationships. Perhaps most importantly, the evaluation methodology must respect the temporal order of the data to avoid data leakage and ensure that the model's performance is assessed in a realistic manner. The model would, for instance, produce over-optimistic results if it has access to future data points when making predictions for a given time step.

2.2.2 Class Imbalance

Perhaps not introduce this yet? It is quite an important topic, though, as it introduces one of the biggest challenges. mFRR up activations are particularly rare events in the NO1 bidding zone compared to down activations and non-activations. This *class imbalance* presents a significant challenge for model training and evaluation. Most ML algorithms are designed to optimize overall accuracy, which can lead to models that are biased towards the majority class, since correctly predicting the majority class contributes more to overall performance. If the minority class is important, as is the case with mFRR up activations, special considerations must be taken to ensure that the model learns to effectively recognize and predict these rare events.

2.2.3 Tree-Based Models

Random Forest / Extra Trees

Gradient Boosting

3 Literature Review

This chapter reviews literature relevant to balancing market forecasting, first outlining the unique characteristics of balancing activation markets and the challenges they present. Then existing methodologies for handling these challenges are presented, before finally identifying specific research gaps that this study and future work can address.

3.1 LOOKING AT PROPOSED ACTIVATION UNCERTAINTY VARIANTS

- **Scenario-Based Activation Models:** Model uncertainty as scenarios. Stochastic optimization. Håberg & Doorman (2017) - reserve dimensioning, TSO perspective [9].
Chance-constrained, imbalance scenario simulations, reserve allocation and dimensioning [10]. Looks at introducing mFRR capacity market participation for hydropower as it was new at the time. Two-stage stochastic participation in day-ahead market, mFRR capacity and energy. Uncertainty in prices modelled with scenarios [11].
- **Activation Ratio or Expected Activation:**

3.2 mFRR Energy Activation Market Characteristics

Balancing markets are by nature unpredictable, as their primary function is to maintain system stability in the face of unforeseen imbalances between supply and demand. This inherent unpredictability is in most balancing markets handled through the capacity market mechanism. The mFRR energy activation market distinguishes itself by only compensating participants for actual energy delivered during activation events. Participants must also bid in the correct direction of activation (upward or downward) to be eligible for activation. When an up-regulating activation is required, the up-regulation price is by design higher than the day-ahead market price, and vice versa for down-regulating activations [12]. These market characteristics make it lucrative for participants to predict activation events accurately, as successful predictions can lead to significant financial gains.

In 2022, Klæboe et al. [12] analyzed day-ahead market bidding strategies for flexible generators taking the balancing power market into account. They found near-zero gains from incorporating balancing market predictions into day-ahead bidding strategies. They discuss, however, that the need for balancing services will increase in the future, and that such strategies will therefore become more relevant and profitable. In Svenska Kraftnät's balancing market outlook 2030 [13] they present that the mFRR capacity demand has and will steadily increase. The report also suggests that since the automated mFRR EAM is only an intermediate step for connecting to MARI (Manually Activated Reserves Initiative), which is an upcoming European-wide mFRR market, further increases in mFRR demand are to be expected. This suggests that predicting mFRR activations will become increasingly important for market participants seeking to optimize their market strategies.

The mFRR energy activation market transitioned from an hourly to a 15-minute resolution as of 4th March 2025, an endeavor aimed at enhancing market efficiency and integrating renewable energy sources more effectively [14]. Under the previous hourly structure, activation signals were constrained to coarse discrete time blocks. Thus, short-lived imbalances or, for instance, rapid ramps in renewable generation could not be reflected optimally in activation decisions. moving to a 15-minute resolution reduces this discretization effect [15].

A study by Kallset and Farahmand found that increased resolution significantly reduces such structural imbalances and achieves about 60% of the possible reduction in total balancing, compared to a 5-minute resolution ideal [15]. Their findings imply that imbalances are now corrected more accurately and efficiently on shorter time scales, making activation patterns more sensitive to rapid system changes. Consequently, the dynamics of the mFRR energy activation market have become

more granular and potentially more volatile, increasing the relevance, but also the difficulty, of short-term activation forecasting.

3.3 Existing Approaches to Balancing Activation Market Forecasting

Forecasting in balancing markets has developed along two main methodological paths. The first focuses on predicting continuous system variables such as imbalance volumes or imbalance prices, mostly appealing to system operators. The second path targets the discrete activation decisions made in the mFRR energy activation market, which are directly relevant for market participants who must place bids in the correct direction. The following sections review key contributions in both areas, highlighting their methodological approaches, limitations, and relevance for short-term activation forecasting under the current 15-minute Nordic market design.

3.3.1 Imbalance Volume and Price Forecasting

Imbalance volume forecasting plays a central role in balancing markets [16] [17]. At the system level, transmission system operators (TSOs) rely on accurate imbalance volume forecasts to minimize balancing costs and ensure system stability. As market participants, accurate imbalance volume predictions are valuable, as this can inform short-term bidding strategies in balancing markets, particularly in markets with high renewable penetration. Consequently, a substantial body of recent research has focused on developing models that forecast system imbalance over short-horizons.

At the system-operator end of the spectrum, several studies develop point forecasting aimed primarily at optimizing TSO operations. One example is the work of Singh et al. [16], who propose a regression-based approach for short-term forecasting of imbalance volumes. They argue that increased renewable variability, combined with the 15-minute activation window, necessitates accurate short-horizon forecasts to enable transmission system operators to anticipate and manage imbalances more effectively. One of their key findings is that their best-performing model reduces the balancing costs by 44.51% compared to TSO-based forecasting. The lower forecast errors lead to reduced costs of energy not supplied, excess energy, and corrections.

While Singh et al. obtain important results for short-term imbalance forecasting in the Belgian power system, related work examines imbalance forecasting in the Nordic region, where renewable mixes and activation patterns may differ. Edling & Azarang (2025) analyze short-term mFRR activation volumes across the four Swedish bidding zones using LSTM-based (Long Short-Term Memory) regression models [18]. Their study highlights strong geographical heterogeneity in activation patterns, identifying structural differences between bidding zones. Price region SE2 exhibited the best forecasting accuracy, while zones SE3 and SE4 showed limited predictability due to high zero-activation frequencies. This finding spotlights an important consideration: regions with frequent imbalance, for instance due to a high share of wind power, do not necessarily have high activation frequencies. SE2 appears to have high activation frequencies due to its significant share of Sweden's flexible hydropower resources, which are often called upon to balance the system. Thus, if an imbalance occurs in SE4, the TSO may choose to activate reserves in SE2 if it is more economical and transmission constraints allow it. Overmaat's study on balancing contributions in the Nordic electricity system, though relatively dated (2019), underscores this point further by concluding that SE1 and SE2 dominate the Swedish balancing contributions on the short and medium time scales [19].

Both Singh et al. and Edling & Azarang aim to provide system operators with actionable insights and tools to better forecast and manage imbalances. While better TSO forecasting does not alter physical imbalance magnitudes, it improves the timing and precision of activation decisions, making activation volumes more closely reflect the real-time system state. For market participants, this suggests that accurate system imbalance forecasts may offer predictive power regarding coming activations. This perspective is developed further in participant-oriented studies such as Plakas et al. [17], who propose a two-stage probabilistic framework for forecasting imbalance volumes and prices sequentially in the Greek balancing market. The first stage employs quantile regression to generate probabilistic forecasts of system imbalances. The second stage leverages the quantiles to

predict imbalance prices. Plakas et al. find that system imbalance volumes are critical predictors of imbalance prices, underscoring the correlation between these two variables. By extending from imbalance volume point forecasts to price forecasting, Plakas et al. provide more actionable insights for market participants seeking to capitalize on opportunities in the balancing market. Their results show that imbalance volumes strongly influence imbalance prices, indicating that system-state indicators provide valuable information for bidders seeking to anticipate balancing-market outcomes.

3.3.2 Activation Forecasting

Whereas imbalance forecasting estimates system imbalance magnitudes, activation forecasting predicts TSO's decision to activate upward, downward, or no mFRR energy. Activation direction is directly relevant for market participants, as balancing market bids must be placed in the correct direction to be eligible for activation. As the Nordic system transitions to a 15-minute market resolution, short-term activation forecasting has become increasingly important, yet the academic literature on this topic remains sparse.

Svedlindh and Yngveson [20] examine the general price formation in intraday and mFRR markets. Among other explorations, they develop logistic regression and ANN (Artificial Neural Network) models to predict activation direction in the mFRR activation market. The ANN model outperforms the logistic regression, achieving solid *accuracy* and *F1-scores*. They identify, however, that *class imbalance* poses a significant challenge, as no-activation events dominate the dataset. This imbalance skews model performance, making it difficult to accurately predict the less frequent upward and downward activations. Despite these acknowledged challenges, Svedlindh and Yngveson achieve promising results. They find that mFRR capacity market prices and procured volumes are informative predictors of activation direction. Since the capacity market is cleared a day ahead with hourly procurement, its outcomes appear to correlate with subsequent activation behaviour.

A recent study in 2025, Porras [21] investigated the use of XGBoost for short-term forecasting of mFRR activation direction and imbalance prices with a two-stage multi-horizon model. This paper is closely related to the present study, as it also focuses on predicting mFRR activations. However, Porras leans heavily into XGBoost, while this study explores a broader range of models. Additionally, Porras produces predictions at an hourly resolution, whereas this study aims for a finer 15-minute resolution - aligning with the recent Nordic mFRR market reform. Porras targets the SE2 price area in Sweden, thereby incorporating data from different sections of the Nordic power system. The dynamics may be similar, but this study focuses exclusively on the NO1 price area in Norway, which may exhibit different characteristics. The primal gap in Porras' study, which this study aims to address, is the opportunity opened by the recent 15-minute market reform, which provides more granular data and potentially improved prediction capabilities.

3.3.3 Literature Synthesis

The reviewed studies differ in their forecasting targets, methodological approaches, and intended users. Regression-based studies such as Singh et al. [16] and Edling & Azarang [18] focus on point forecasts of system imbalance volumes. These models primarily serve system operators and rely on data streams that are not always available to market participants. While such point forecasts provide valuable information about the magnitude of expected imbalances, they do not directly predict activation direction and do not quantify uncertainty.

Probabilistic approaches, exemplified by Plakas et al. [17], attempt to characterize uncertainty explicitly by producing distributions over imbalance volumes and prices. This is more directly useful for market participants, but their framework till does not address the discrete up/down/none decision that governs mFRR activations.

Only a small number of studies examine activation direction itself. Svedlindh and Yngveson [20] develop logistic regression and ANN classifiers, showing that activation direction is in principle predictable but heavily constrained by class imbalance and the use of capacity-market features that

are not available in real-time to market participants. Porras [21] investigates activation direction using XGBoost at an hourly resolution in SE2, demonstrating the potential of machine-learning classification but leaving open questions regarding finer temporal resolution and the feasibility of predictions outside high-activation regions.

Taken together, the literature indicates that imbalance information carries predictive value for prices and activations, but existing approaches either (i) focus on continuous volume forecasting, (ii) rely on non-participant data, or (iii) operate at coarser temporal scales than current Nordic market design. There remains limited evidence on whether activation direction can be predicted reliably using only features available to market participants and at the new 15-minute resolution.

Table 1: Comparison of key studies on balancing market forecasting

Study	Target Variable	Resolution	Model Type
Singh et al. (2025)	Imbalance volume	15-min	Regression
Edling & Azarang (2025)	mFRR activation volume	1-hour	LSTM
Plakas et al. (2025)	Imbalance volume + price	1-hour	Probabilistic / QR
Svedlindh & Yngveson (2025)	Activation direction	1-hour	Logistic / ANN
Porras (2025)	Activation direction + price	1-hour	XGBoost

Table 1 summarizes the key characteristics of the reviewed studies, reviewed studies, highlighting their differing targets, resolutions, and model types.

3.3.4 Research Gap

Despite notable progress in balancing-market forecasting, several gaps remain in the existing literature. Most existing studies focus on continuous imbalance or activation volume forecasting at an hourly resolution. While these approaches provide valuable insights for system operators, they do not address the discrete up/down/none activation decisions faced by market participants, especially under the newer 15-minute market design.

Studies that do consider activation direction (Svedlindh & Yngveson and Porras) do so at an hourly resolution. With the transition to 15-minute activation intervals, however, finer temporal patterns and short-lived imbalances may become more relevant. Investigating whether higher time resolution can improve prediction performance is therefore of interest.

Geographically, relatively little work has focused on the Norwegian bidding zones. The reviewed studies investigate the Belgian and Greek systems (Singh et al. and Plakas et al.) or Swedish bidding zones (Edling & Azarang, Svedlindh & Yngveson, and Porras). Swedish bidding zones naturally share similarities with Norwegian zones, as both are part of the Nordic system. However, differences in generation mixes, cross-zonal flows, and local balancing tendencies may lead to distinct activation patterns.

A further consideration is that few studies robustly address the data availability constraints faced by market participants. Some studies incorporate TSO-only data, while others do not explicitly evaluate whether their chosen features would be accessible to market participants at decision time. As a result, the extent to which activation direction can be predicted under participant-feasible information constraints remains relatively open.

Taken together, these observations suggest the value of a systematic evaluation of activation-direction forecasting in the NO1 bidding zone, using features available to market participants and operating at the updated 15-minute resolution. This study contributes to this area by comparing a range of machine-learning models and assessing their performance under realistic data and operational constraints.

4 Methodology

4.1 Overview of Methodological Approach

The methodological approach adopted in this project is visualized in Figure 4. The first step, data collection, involves gathering relevant datasets from various sources, including Nord Pool, NUCS, and ENTSO-E. The collected data is then preprocessed and cleaned in the second step. These first steps are uniquely colored red in the figure to indicate that they are mostly one-time efforts required to set up the dataset for further.

The third step, feature engineering, involves creating and selecting relevant (**either introduce features here or make sure it has been introduced**) features, i.e. attributes in the dataset that help the model learn patterns related to mFRR activations. The feature-engineered dataset is subsequently sent to the model training algorithm. Here, machine learning techniques are applied to train classification models using the prepared dataset. Then, the trained models are evaluated with the relevant metrics to assess their predictive performance. These three steps are colored purple, indicating that they are iterative processes that make up the bulk of the work invested in this project.

Evaluation results are interpreted in the last step indicated by the blue box. This step is unique as it only involves deriving insights and conclusions from the previous steps. Methodological feedback loops flow from the interpretation step back to the feature engineering, model training, and evaluation steps. These feedback loops represent the iterative nature of the methodology, where insights gained from interpretation inform further refinements and improvements in the earlier stages of the process.

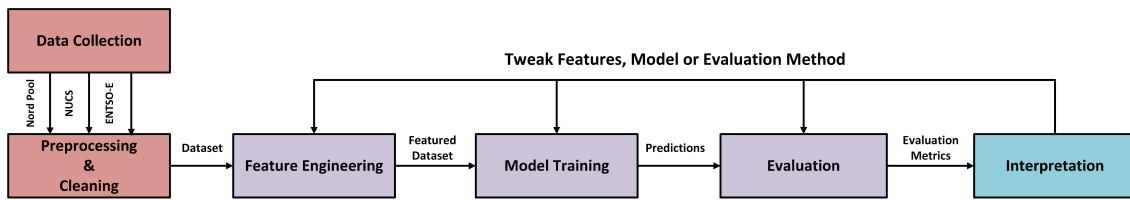


Figure 4: Overview of the methodological approach used in this project.

4.2 Data Sources

This section introduces the data sources used in this project and describes how the data was gathered and preprocessed for further use in model training and evaluation. The gathered data spans a period from January 1, 2024, to December 4, 2025, providing a comprehensive view of NO1 mFRR activation patterns over nearly two years.

4.2.1 Nord Pool

Nord Pool provides market and system data in the Nordic region. Data was downloaded manually through the Nord Pool data portal in yearly chunks [22]. A manual approach was necessary due to the Nord Pool API not being available for this project. Data API access requires a commercial agreement with Nord Pool, which was not obtained. This is not an issue for this project and research as real-time data is not required for model training and evaluation. If the models were to be deployed in a real-time setting, however, access to real-time data through the API would be essential. The following subsections describe the specific datasets obtained from Nord Pool. **This is not really relevant right here, but it is an important point to make somewhere.**

mFRR activation data. The primary dataset used in this study consists of manual Frequency Restoration Reserves (mFRR) activation data from the Nordic electricity market, specifically for the bidding zone NO1. This data includes accepted and activated up- and down-regulation bids at a 15-minute resolution. The activated volumes provide the target variable for the prediction models, indicating whether an mFRR activation occurred in a given 15-minute interval.

Cross-zonal flows. Cross-zonal flows refer to the electricity flows between different bidding zones in the Nordic market. In this project, only cross-zonal flows involving the NO1 bidding zone are considered: flows between NO1 and SE3, NO1 and NO3, NO1 and NO5, and NO1 and SE3. These flows are crucial for maintaining grid stability and optimizing the use of available resources. The dataset includes information on cross-zonal flows to provide additional context for mFRR activations.

Load and production data. Load and production forecasts provide insights into the expected system state. Forecast may on their own provide valuable information about potential mFRR activations, but when combined with actual load and production data, the model can learn to identify discrepancies between expected and actual system states. Such discrepancies often lead to imbalances that require mFRR activations to restore balance. The different production sources (e.g., hydro, wind, thermal) have varying characteristics and impacts on grid stability. Among them, wind power is particularly relevant due to its intermittent nature, which can lead to sudden changes in generation levels. Wind power production data is therefore predicted to have the biggest impact on mFRR activations among the different production types.

Load/consumption data is much simpler in nature, as *who* or *what*, essentially the source of consumption, is not as relevant as the source of production. Consumption forecasts and actual consumption data can still be useful, however, as sudden changes in consumption patterns can lead to imbalances that require mFRR activations.

4.2.2 NUCS

NUCS, or the Nordic Unavailability Collection System, is a service for collection of data on unavailable data in the Nordic power system. NUCS is an important part of this project, as it provides otherwise unavailable data that served as features in the models. NUCS is unique from the other data sources used in this project, as it provides data through an API (Application Programming Interface) [23]. This allows for automated data retrieval, which is especially useful for real-time applications. This project does not have access to comprehensive real-time data, and the NUCS API was thus only used to gather historical data for the training and evaluation of the models. An algorithm was developed, however, to automatically retrieve up-to-date data from the NUCS API for potential future real-time applications.

aFRR data. aFRR data is not available through Nord Pool, but through the NUCS API, historical aFRR procurement prices and volumes for the NO1 bidding zone are accessible. The data is available at an hourly resolution, with separate values for up- and down-regulation. This data provides insights into the amount of balancing that is expected to be needed in the system.

4.2.3 ENTSO-E

ENTSO-E, the European Network of Transmission System Operators for Electricity, is a key organization in the European electricity market. ENTSO-E provides a wide range of data related to electricity generation, consumption, and grid operations across Europe [24].

aFRR Activation Data aFRR prices and capacities are available through NUCS as discussed earlier, but aFRR activation data is not. However, ENTSO-E provides detailed data on aFRR

regulations via their "Accepted Offers and Activated Balancing Reserves" dataset. This dataset is supposed to provide information about up and down regulations from all balancing markets. Only aFRR data seems to be available, however, which is sufficient for this project. The data is available at an hourly resolution, which is coarser than the 15-minute resolution of the mFRR data. This data is, in the same manner as other 1-hour resolution data, resampled with forward filling to match the 15-minute resolution of the main dataset.

4.3 Data Preprocessing and Analysis

4.3.1 Dataset structure

The data is represented as a time series, where each record in the dataset consists of a set of attributes connected to one point in time. More specifically, the data contains a sequence of 15-minute interval time stamps. Each time stamp may or may not have an associated activation, which is the target variable the model is trying to predict. The features describe the system state at that time stamp, providing context for the model to learn from.

4.3.2 Resampling, Imputation, and Merging

Nordic power market data is transitioning from hourly to 15-minute resolution. However, many datasets are still only available at an hourly resolution, and some datasets have mixed resolutions over the past years. Day ahead price data, for instance, transitioned to 15-minute market time units (MTU) on September 30th 2025 [25]. Consequently, data before this date is at an hourly resolution, while data after this date is at a 15-minute resolution. mFRR activation data transitioned to 15-minute MTU on March 4th 2025 [3], but Nord Pool has updated their historical data to be at a 15-minute resolution for the entire dataset period.

To ensure consistency across all datasets, all data with hourly resolution data subsets are resampled to a 15-minute resolution. This is done using the built-in Pandas `resample()` function with forward filling. Forward filling entails propagating the last valid observation forward to fill gaps. For instance, if the day-ahead price between 10:00 and 11:00 is 50 EUR/MWh, then after resampling, the price for 10:00, 10:15, 10:30, and 10:45 will all be set to 50 EUR/MWh, as illustrated in Figure 5. Thus, the data still remains constant within the hour, but is now available at the desired 15-minute resolution.

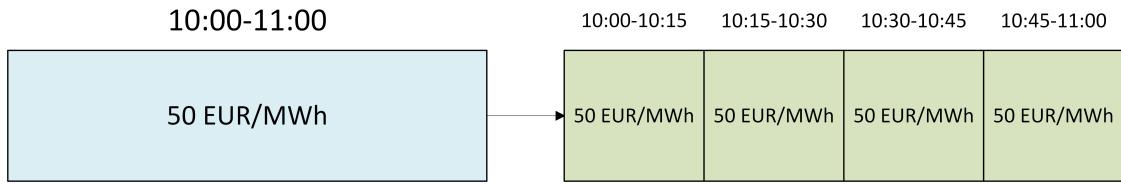


Figure 5: Visualization of resampling from 1-hour to 15-minute resolution using forward filling.

The handling, or *imputation*, of missing values, is an important step in data preprocessing. In time-dependent data, simply removing the rows with missing values is problematic, as it would break the time continuity. Most of the datasets used in this project do not have significant issues with missing values, but for the few that do, interpolation or forward/backward filling techniques are used to estimate the missing values based on surrounding data points. Interpolation is often preferred, as it can provide smoother estimates, as it considers both previous and subsequent data points. Backward filling is used when missing values are at the beginning of the dataset, as there are no previous data points to reference. Otherwise, forward filling is used as the default method, as it maintains the most recent known value, thus preventing time leakage from future data points.

After resampling and handling missing values, the various datasets are merged into a single dataset.

A successful merge requires that all datasets share a common time index format and resolution. Datasets from different sources often differ in time zone and how time stamps are encoded. Therefore, all time stamps are converted to a common time zone (CET/CEST) and format (Pandas `to_datetime()` function) before merging. The final merged dataset contains all data aligned at a 15-minute resolution, ready for feature engineering and model training.

4.3.3 Time Restrictions

Short-term prediction is heavily reliant on the availability of recent data.

aFRR activation data is naturally not available at the target time $t + 4$ and prediction time t as these activations has not yet occurred or finished. The 1-hour resolution is problematic as it essentially causes the data to become available in clusters of four 15-minute intervals at a time. This causes inconsistencies in the what specific lag features are available at different prediction times. For instance, assume data becomes available immediately after the hour. Then, at time $t = 10:00$, the last available aFRR activation data would be from 09:00-10:00. This means that lag features for $t - 1$, $t - 2$, $t - 3$, and $t - 4$ would be available. However, in the worst case, at a prediction time $t = 10:45$, the last available aFRR activation data would still be from 09:00-10:00. This means that only lag features starting at $t - 4$ would be available, while $t - 1$, $t - 2$, and $t - 3$ would not. This inconsistency in available lag features complicates the model training and evaluation process. There are multiple ways to handle this issue:

- **Option 1:** Only use lag features that are always available, i.e., only use lag features starting from $t - 4$ and further back. This is the easiest option to implement but it sacrifices potentially valuable information from more recent time steps.
- **Option 2:** Create separate models for different prediction times within the hour (e.g., one model for predictions at 00, 15, 30, and 45 minutes past the hour). Each model would be trained with the each with their own set of available lag features. This approach maximizes the use of available data but requires maintaining multiple models. The more gathered datasets are of a 1-hour resolution, the more valuable this approach becomes.
- **Option 3:** Use imputation techniques to estimate the missing lag features based on available data. This approach allows for a single model to be used but introduces uncertainty due to the imputation process.

Restrictions In an ideal world, a market participant would be able to choose to bid or not in real time. The mFRR activation market, however, is not ideal in this sense. Bids are accepted for time slots $t + 4$ intervals into the future, i.e. one hour into the future. This greatly restricts market actors in terms of how fast they can act on recent information. In fact, at the time of bid closing for a time slot $t + 4$, only mFRR activation data from $t - 3$ and earlier is available. Thus there are seven 15-minute intervals of unavailable data. This data would be the most useful for the model, since activation data at time t correlates with the system state and activations close to t .

It would, for instance, be considerably easier to identify streaks of activations, as activations often occur in long contiguous sequences. In the real world scenario, data from $t - 3$ and before is still useful, but can often be misleading and cannot be trusted in isolation. The dataset developed in this project investigates ways to provide the model with crucial context about the system's state of stress. The intuition is that data such as physical cross-zonal flows in and out of NO1 may correlate with NO1 mFRR activations. When many such features are combined, the model may develop complex relationships between them, enhancing its predictive performance.

Most of the available and useful data are not available in real time. For instance, consumption and production data are often published with a delay of several hours. The model must therefore rely on features that are available in real time, or with minimal delay. Such features often take the form of forecasts, which are available ahead of time. Forecasts are inherently uncertain, but they still provide valuable information about the expected system state.

4.3.4 Exploratory Data Analysis

Target Class Imbalance. Figure 6 displays the distribution of mFRR activations over the dataset period. The figure especially highlights the infrequent nature of up-activations, which occur far less often than down-activations. This phenomenon induces a *class imbalance* in the prediction task, which in general makes it more challenging for models to accurately predict the minority class [26].

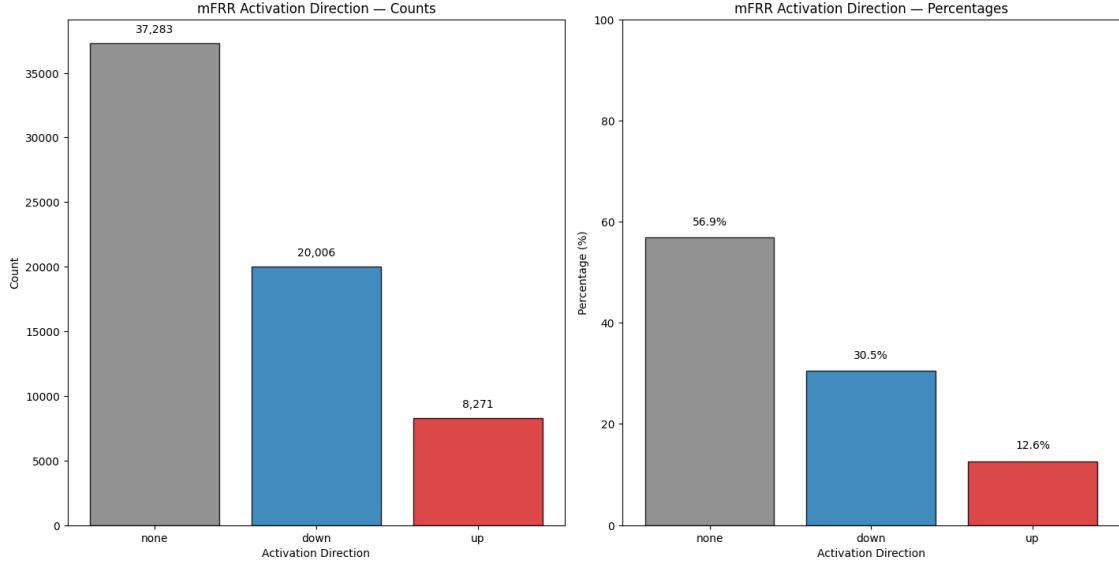


Figure 6: mFRR activation distribution.

The distribution also reveals the majority class with 56,9% of all intervals having no activation. Activations occur in 43,1% of the intervals, with down-activations being the most common at 32,5% and up-activations being the least common at 10,6%. Thus, if one were to consider the binary case of activation vs. no activation, the classes would be approximately balanced. Distinguishing between up- and down-activations, however, makes the problem more nuanced and challenging.

Temporal Activation Patterns

Feature Distributions. Figure 7 shows the distribution of cross-zonal flow directions for the NO1 bidding zone. The figure indicates that flows between NO1 and NO3, NO5, and SE3 are drastically skewed towards imports into NO1, while flows between NO1 and NO2 are reversely skewed. NO1-NO3 and NO1-NO5 show the most pronounced skewness, with very few occurrences of exports from NO1 to these zones. **This might be interesting to talk about later, as imbalance here may make the directionality more predictive.**

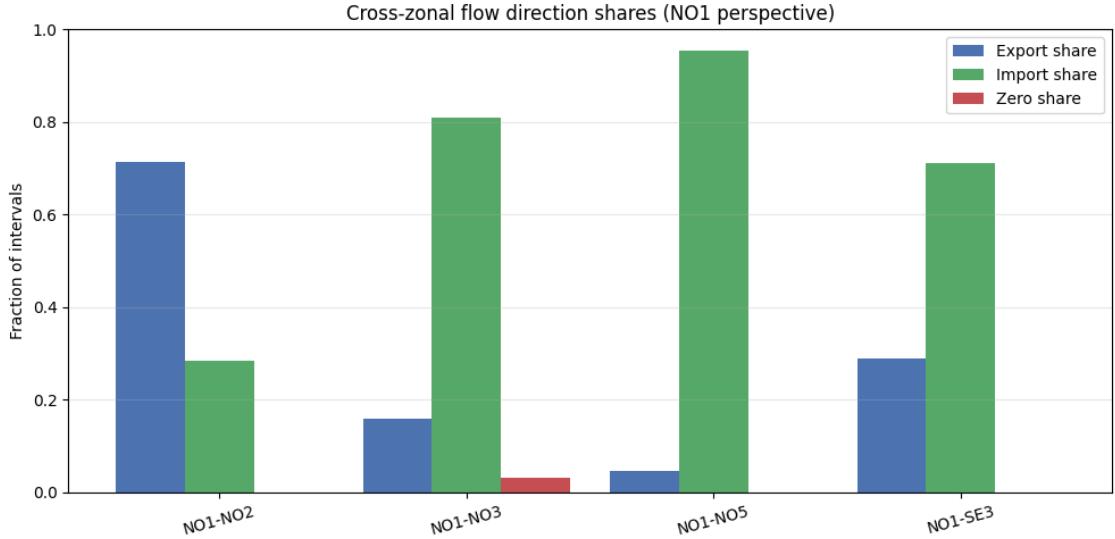


Figure 7: Cross-zonal flow distributions for the NO1 bidding zone.

Figure 8a and 8b show the relative utilization of the cross-zonal connections for NO1 as density plots. The figures illustrate how heavily the connections are utilized for imports and exports, respectively. The utilization is calculated as the ratio between actual flow and the NTC (Net Transfer Capacity) capacity of the connection. The export utilization figure highlights the lack of exports from NO1, except for the NO2 connection, indicated by the tail thickness between 0.4 and 1.0 utilization. The import utilization figure, on the other hand, shows that all connections, except NO2-NO1, are heavily utilized for imports, with many thick tails approaching full utilization. The NO1-NO2 connection is thus almost exclusively used for exports from NO1 to NO2, whilst the other connections are primarily used for imports into NO1.

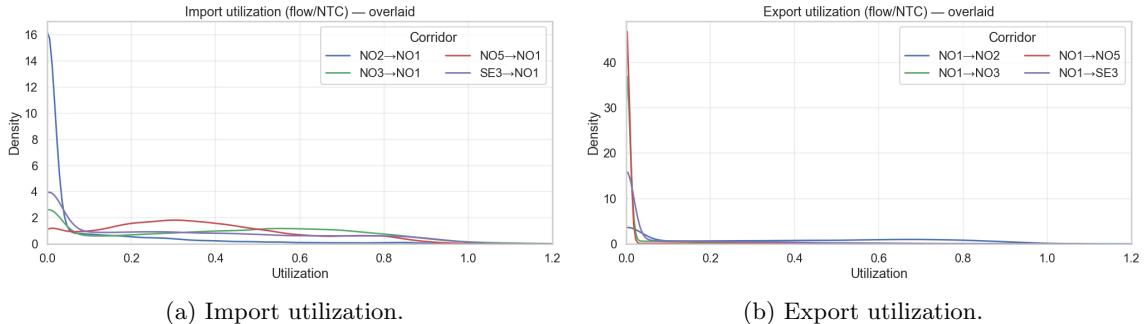


Figure 8: Cross-zonal flow utilizations for the NO1 bidding zone calculated as the ratio between actual flow and NTC capacity.

Production

Table 2: Summary statistics for wind-related features (2024–2025, NO1)

Metric	Mean	Std	Min	P10	P50	P90	Max	Count
Wind DA Forecast	121.64	96.34	0.0	16.0	95.0	274.0	370.0	62,680
Wind Intraday Forecast	135.80	104.55	0.0	15.0	111.0	294.0	376.0	38,972
Wind Actual Production	120.38	102.72	0.0	8.0	91.0	283.0	380.0	65,568
Wind Revision (ID-DA)	12.59	15.27	0.0	1.0	7.0	30.0	151.0	38,972
DA-Actual Error	0.03	38.85	-250.0	-44.0	-2.0	47.0	221.0	62,680
ID-Actual Error	2.97	35.30	-227.0	-38.0	1.0	46.0	189.7	38,972
Abs DA Error (%)	0.58	0.99	0.0	0.038	0.24	1.36	5.0	61,486
Abs ID Error (%)	0.42	0.75	0.0	0.032	0.19	0.92	5.0	38,370
Wind Share	0.054	0.047	0.0	0.0036	0.040	0.128	0.228	65,568

aFRR. Figure 9 and 10 show a histogram and time series plot of the hourly NO1 aFRR procurement prices between January 1, 2024, and 1. December 2025. Both figures highlight a problem in pre-July 2024 data, as this period contains many missing values. The problem seems to be intermittent missing values rather than large gaps of missing data, indicated by the time series plot. Interpolation is thus a suitable imputation method, as it can estimate the missing values based on surrounding data points. This may miss out on some extreme price spikes, but is still expected to provide a reasonable estimate. The data appears complete after this date. **Short discussion on price levels and patterns here would be good. ALSO be specific on up/down.**

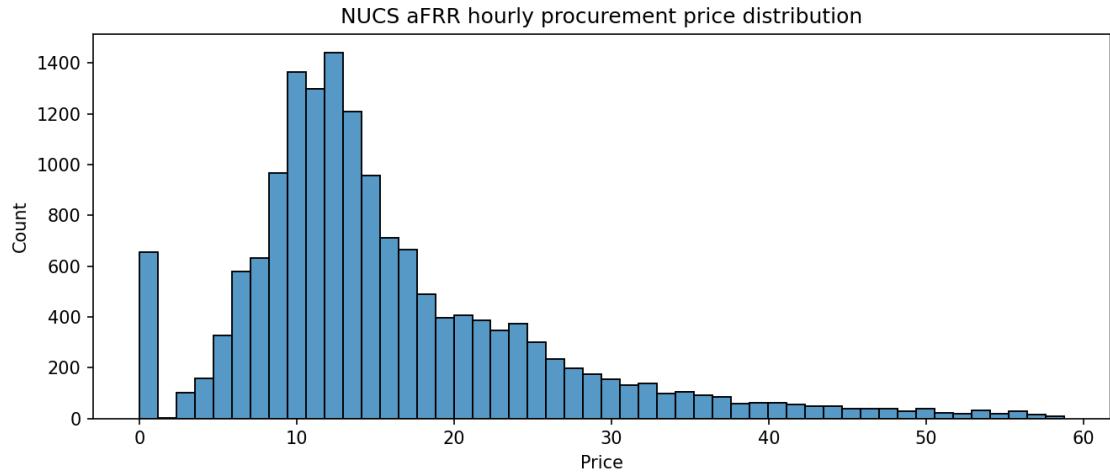


Figure 9: A histogram of hourly aFRR procurement prices for the NO1 bidding zone from NUCS data.

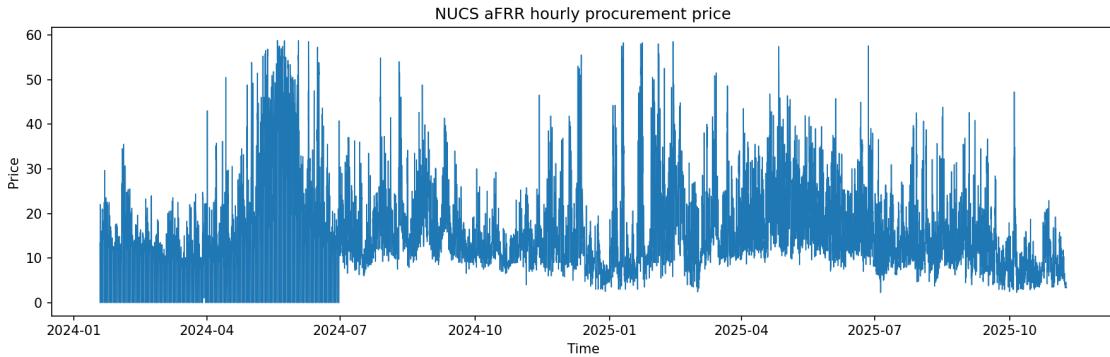


Figure 10: Hourly aFRR procurement prices for the NO1 bidding zone from NUCS data.

Correlation Analysis

Information Availability

4.4 Feature Engineering

Features are attributes in a dataset that describe each data point. A dataset for predicting a person's income could, for instance, have features like gender, job type, and age. Then, each data point represents a person and relevant information about the person in terms of the target variable – income. In this project, each data point represents a specific time stamp in the mFRR activation dataset, and the features describe the system state at that time. This includes information such as electricity demand, generation capacity, and market prices, all of which can influence mFRR activations.

Already available features can be transformed to create new higher-level features that may better capture the underlying patterns in the data. For example, if one has features for year of death and year of birth, a new feature for age at death can be created by subtracting the year of birth from the year of death. This is known as *feature construction* [27]. Features can be constructed in various ways, such as through mathematical operations or aggregations. This project leverages this concept to create new features that may enhance the model's predictive capabilities.

Feature selection is crucial. There are many features that may seem useful and relevant in isolation, but sometimes they mislead the models, or they work poorly in combination with other seemingly good features. Theoretical analysis of the usefulness of certain features can be helpful, but only trial-and-error together with feature-importance analysis will uncover the features' actual impact on performance. Feature selection will be subject to restrictions outlined in Section 4.3.3 to ensure that only real-time available features are used.

4.4.1 Lag features

A lag feature is a feature that represents the value of a variable at a previous time step. Lag features are commonly used in time series analysis to capture *temporal* dependencies and trends in the data [28]. By including lag features, the model can leverage historical information to make more informed predictions about future mFRR activations. For instance, if there was an upregulating activation in the previous time step, it may indicate a higher likelihood of another upregulating activation in the current time step.

Activation lag features Activation lag features are the most important lag features for this problem, as they convey important information about recent temporal activation trends. They are, however, restricted by the real-time limitations, so the model may only use activation lag features from $t - 3$ and earlier for predicting activations at time $t + 4$. As a result, lag features for upregulating and downregulating activations are created for time steps $t - 3, t - 4, t - 5, \dots, t - 9$. These features are useful by themselves, but they also serve as a basis for creating other features that capture activation trends more effectively.

Persistency (streak length) Let t denote the reference time such that the model predicts activation at $t + 4$. At bid close, only activation data up to and including $t - 3$ are available. Define binary activation indicators $A^\uparrow(\cdot)$ and $A^\downarrow(\cdot)$ for up- and down-regulating activations, respectively. The up- and down-persistency features are the lengths of the most recent consecutive activation runs ending at $t - 3$: **TODO: Consider only formulating persistency with words, as the mathematical expression is more complex than the concept.**

$$S^\uparrow(t) = \sum_{j=0}^{\infty} \prod_{i=0}^j A^\uparrow(t - 3 - i), \quad S^\downarrow(t) = \sum_{j=0}^{\infty} \prod_{i=0}^j A^\downarrow(t - 3 - i).$$

Here, j is the (zero-based) window length and i is the offset index within that window. The product is 1 only while all the last $j+1$ values are 1, so when a 0 is encountered, the product becomes 0 for all larger j . Thus, the sum counts how many consecutive intervals back from $t - 3$ remain all ones before the first zero. In words, $S^\uparrow(t)$ (resp. $S^\downarrow(t)$) counts how many consecutive up- (resp. down-) activations occurred immediately before $t - 3$. If there was no activation at $t - 3$, the corresponding streak length is 0.

Design choice: separate vs. signed. When designing the persistency features, a choice had to be made between using two separate non-negative integer features or a single signed feature (e.g., $S^\uparrow(t) - S^\downarrow(t)$). Separate features avoid conflating direction with magnitude and let the model learn asymmetric effects. This option may be easier for the model to interpret, as it provides a clearer distinction between up- and down-regulating activations. The signed variant is more compact, but it may introduce ambiguity in how the model interprets the values. It is, however, important to consider the possibility of persistency overreliance. If the model relies too heavily on persistency features, it may overlook other important factors influencing activations. This could lead to suboptimal predictions, especially in scenarios where activation patterns change. Only having one persistency feature may reduce this risk, as persistency is encoded in a single column. Both options will be explored and evaluated during model development, and the pros and cons of each approach will be assessed based on empirical performance.

Importance. Persistency features are perhaps the single most important features in the dataset. They provide compact and direct information about recent activation trends, making it easier for the model to identify activation streaks. As activations often come in long contiguous runs, and even though there are real-time limitations, the persistency features still provide the model with valuable context. The longer the streak, the more likely it is that the activation trend will continue. In some cases, the model might be able to predict an activation solely based on a high persistency value, without needing to consider other features.

Much of the motivation behind this project is to develop models that can predict activations based on system state features rather than just relying on past activation trends. The reality is, however, that persistency features are extremely dominant, and they may overshadow the contributions of other features.

4.4.2 Cross-zonal flow features

Cross-zonal flow features capture information about electricity flows between different zones or regions in the power grid - in this case, in and out of the NO1 bidding zone. These flows can indicate the grid's stress level and influence mFRR activations. For instance, high inflows into NO1 may signal increased demand or generation shortages, potentially causing upregulating activations. Conversely, high outflows may indicate surplus generation, potentially causing downregulating activations. It is unlikely, though, that cross-zonal flow features alone can predict activations. Combining them with available transfer capacity provides a picture of how close the grid is to its operational limits. For instance, if the inflow into NO1 is close to the maximum available transfer capacity, only a small margin remains for additional inflows, which could increase the likelihood of upregulating activations. Such situations often occur in zones that are short, i.e. zones where consumption exceeds production. In such cases, the grid operator may need to activate expensive mFRR reserves to maintain grid stability when no more cheap imports are possible. It is important that the models developed in this project are able to capture these kind of relationships as they are among the most valuable for a potential user of the models.

Capacity-normalized cross-zonal flow. Raw cross-zonal flow magnitudes are not comparable across interconnections or over time because each line has different capacity and the available transfer capacity (ATC) varies. The same absolute flow can be insignificant on a strong interconnection but critical on a constrained one. To obtain a dimensionless, capacity-normalized measure

of proximity to operational limits—and to make features comparable across borders and time—flows are expressed as a ratio to the relevant directional ATC.

Let $F_i(t)$ be the *signed* flow for interconnection i at time t (positive into NO1, negative out of NO1). Let $ATC_i(t) \geq 0$ denote the available transfer capacity magnitude used for normalization (e.g., a symmetric ATC for interconnection i at time t). The capacity-normalized ratio is then

$$F_{\text{ratio}}^i(t) = \frac{F_i(t)}{ATC_i(t)}.$$

Hence $F_{\text{ratio}}^i(t) \in [-1, 1]$ when flows are within limits, approaching +1 as inflow nears the capacity and -1 as outflow nears the capacity.

4.4.3 Temporal features

Temporal features capture time-related patterns in the data. These features help the model understand how mFRR activations vary with time, such as daily or weekly cycles. Basic temporal features include hour of the day, day of the week, and month of the year. These features allow the model to learn patterns related to specific times. mFRR activations could, for instance, be caused by completely different factors during peak hours on weekdays compared to off-peak hours on weekends. Temporal features like these are most often represented using cyclical encoding to reflect their periodic nature. For example, 1 AM and 11 PM are close in time, even though their numerical representations (1 and 23) are far apart. Cyclical encoding uses sine and cosine transformations to capture this periodicity [29]. Hourly features are, for instance, encoded as:

$$\text{Hour}_{\text{sin}} = \sin\left(2\pi \cdot \frac{\text{Hour}}{24}\right), \quad \text{Hour}_{\text{cos}} = \cos\left(2\pi \cdot \frac{\text{Hour}}{24}\right).$$

Monthly features can be encoded similarly, using 12 as the divisor instead of 24.

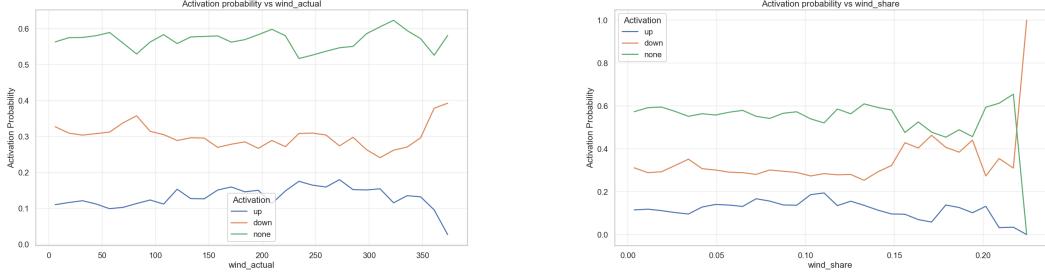
4.4.4 Price features

Price features capture information about the various electricity market prices. The mFRR activation market is closely linked to other electricity markets, such as the day-ahead market, the intraday market, and the aFRR market. Most prices may not have direct impacts on activations, but by crafting features that capture important relationships between prices, the model may be able to infer system stress levels that could lead to mFRR activations. Large discrepancies between day-ahead prices and intraday prices may, for instance, indicate unexpected changes in supply or demand, which should correlate with mFRR activations. Similarly, the difference between aFRR prices and mFRR prices may provide insights into the relative costs of balancing services, which could influence activation decisions.

4.4.5 Load features

4.4.6 Production features

Production features capture information about electricity generation, particularly from renewable sources like wind power. Wind power production features were considered promising candidates for predicting mFRR activations, as wind power is intermittent and can cause sudden changes in supply. Figures 11a and 11b show values of realized wind production and wind share (wind production as a fraction of total production) plotted against the distribution of mFRR activations. These figures indicate that there is no direct correlation between wind production and mFRR activations. The existence of such a correlation would have made it easy for the model to leverage wind production features for predicting activations. The hope is, however, that wind production features will prove useful when combined with other features, as the model captures complex relationships between features.



(a) Realized wind production plotted against mFRR activation distribution. (b) Forecasted wind production plotted against mFRR activation distribution.

Figure 11: Load and production data distributions.

4.4.7 Interaction features

Interaction features are created by combining two or more existing features to capture complex relationships that may influence mFRR activations.

4.5 Evaluation Framework

Classification problems are often evaluated using accuracy, precision, recall, and F1-score. These metrics are defined as follows [30]:

- **Accuracy:** The ratio of correctly predicted observations to the total observations. It is calculated as:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP is true positives, TN is true negatives, FP is false positives, and FN is false negatives.

- **Precision:** The ratio of correctly predicted positive observations to the total predicted positive observations. It is calculated as:

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall:** The ratio of correctly predicted positive observations to all observations in the actual class. It is calculated as:

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **F1-score:** The weighted average of Precision and Recall. It is calculated as:

$$\text{F1-score} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

While accuracy is a commonly used metric, it can be misleading in cases of imbalanced datasets. For instance, if only 5% of the data points belong to the positive class, a model that always predicts the negative class would achieve 95% accuracy, but would be useless for identifying positive cases. In such scenarios, precision, recall, and F1-score provide a more nuanced evaluation of model performance, especially in applications where the costs of false positives and false negatives differ significantly. In the context of mFRR activation prediction, false negatives (failing to predict an activation) may lead to missed opportunities for market participation, while false positives (predicting an activation when there isn't one) could result in unnecessary costs or penalties. Therefore, a balanced consideration of these metrics is essential for developing an effective classification model.

Multiclass evaluation. When the classification is no longer binary, but multiclass (three or more classes) which is the focus of this project, evaluation metrics must be adapted. Common approaches include macro-averaging, micro-averaging, and weighted averaging. The weighted-averaged F1-score is calculated by taking the mean of all per-class F1-scores, weighted by their *support*, i.e., the number of true instances for each class. This approach is appropriate if the class distribution is imbalanced, and it is desired to give more importance to the performance on the more frequent classes. This is not the case in this project, as minority classes are of great importance. Therefore, the macro-averaged F1-score is used, which treats all classes equally regardless of their frequency. Macro-averaged F1-scores are calculated by computing the F1-score for each class independently and then taking the average of these scores. Micro-averaging, which aggregates the contributions of all classes to compute the average metric, is less commonly used for imbalanced datasets. Micro-averaging is analogous to accuracy in binary classification, which is not suitable for this problem due to class imbalance [31].

Confusion matrix. The confusion matrix is a cross tabulation of predicted versus actual class labels. It provides a breakdown of correct and incorrect predictions for each class, aligning correct predictions along the diagonal. What stands out in a confusion matrix is not just the overall accuracy, but also the specific types of errors the model makes. For instance, in a three-class classification problem, the confusion matrix can reveal if the model tends to confuse certain classes more than others. An example confusion matrix for a three-class classification problem is shown in figure 12. In this example, the model performs well on the 'down' class and 'none' class, capturing most of the true instances while making few false predictions. The 'up' class has four correct predictions but also three misclassifications as 'none' and one as 'down', indicating that the model struggles to differentiate the 'up' class from the others.

		Predicted Classification			
		Classes	Down	None	Up
Actual Classification	Down	10	2	1	
	None	3	12	3	
	Up	1	1	4	

Figure 12: Example of a confusion matrix for a three-class classification problem.

Precision-recall trade-off. This project primarily focuses on maximizing the F1-score, as it balances precision and recall, providing a comprehensive measure of the model's performance in predicting mFRR activations. Accuracy is essentially neglected due to the imbalanced nature of the dataset. Between precision and recall, recall is slightly prioritized, as missing an activation prediction is considered more detrimental than a false alarm in this context. Precision can also be optimized more easily as the confidence threshold can be adjusted post-training to favor precision over recall or vice versa. To achieve a high recall score, the model must be capable of identifying as many actual activations as possible, even if it means occasionally predicting an activation when there isn't one. The model must be gutsy and attempt to identify patterns that indicate an upcoming activation, not just follow recent activation history. This approach aims to ensure that the model has practical utility in real-world market participation scenarios. If the model had no such pattern recognition capabilities, it would be of little use beyond simple statistical analysis of recent activation trends, which could be performed without machine learning.

There is, however, a limit to how much recall can be prioritized. If the model predicts an activation for most time intervals, it will achieve a high recall but at the cost of precision, rendering it ineffective. Therefore, the model must strike a balance, ensuring that it is both sensitive to actual activations and specific enough to avoid excessive false positives. This balance is crucial for the model's success in practical applications, where both types of errors have significant implications. The exact precision-recall trade-off can be adjusted based on the specific use-case. Three potential cases are outlined below:

- **Case 1 - Recall-focused:** A recall-focused approach can be useful for analysis purposes, where the goal is to identify periods of increased risk for activations. In this case, the model can be optimized to achieve a high recall score, even if it means sacrificing precision. Such a model would be valuable for understanding the conditions that lead to activations, but may not be suitable for direct market participation due to the high number of false positives. This approach might even be the best for market participation as the ability to discover more activations could outweigh the costs of false positives.
- **Case 2 - Balanced approach:** For general applications, maintaining a balance between precision and recall is often desirable. The model can be optimized to achieve a high F1-score, ensuring that both metrics are adequately addressed. This involves fine-tuning the model's parameters and threshold settings to find an optimal trade-off. Ideally, this approach would be used, achieving good performance in both precision and recall, making the model versatile for various applications, including market participation. Achieving such results is, however, quite challenging as either precision or recall often needs to be sacrificed to some extent to improve the other.
- **Case 3 - High precision focus:** In situations where false positives carry significant costs, the model can be adjusted to prioritize precision. This may involve raising the confidence threshold for predicting an activation, reducing false positives but potentially missing some true activations. For multi-market actors, such a model is useful, as they can afford to be selective about which market to participate in, only bidding into activation markets when the model is very certain of an upcoming activation.

In this project, the focus is primarily on Case 1 and Case 2, with an emphasis on achieving a high F1-score while slightly prioritizing recall. This approach aims to ensure that the model has a chance to give users novel insights into mFRR activations, potentially providing a competitive edge in market participation.

The transition metric is one more evaluation metric which is important to consider for this specific problem: how often the model manages to predict the start of an activation streak. This is measured by looking at all sequential time interval pairs where the first interval did not have an activation, but the second one did. If the model predicted an activation for the second interval, it is counted as a successful prediction *transition*, which is how this metric will be referred to. This metric is important because it concretizes the model's ability to catch the onset of activation periods, which is the most valuable aspect for market participants. If the model can often enough predict these transitions, it can provide significant strategic value, even if its overall precision and recall are not perfect.

Probability correctness. Although the models make hard classifications, they do so based on predicted probabilities for each class. It is, therefore, useful to evaluate how well these predicted probabilities align with actual outcomes. For example, if the model predicts an activation with a probability of 0.8 and an activation does not occur, this should be considered a more severe error than if the model falsely predicted an activation when the predicted probabilities were more evenly distributed.

4.5.1 Data Splitting

Proper data splitting is crucial for evaluating the performance of machine learning models. In this study, the dataset is divided into training, validation, and test sets based on temporal order to prevent data leakage and ensure that the model is evaluated on unseen data. The training set is used to train the model, the validation set is used for internal hyperparameter tuning and model evaluation during development, while the test set is reserved for the final evaluation of the model's performance.

[Explain more.](#)

4.5.2 Model Selection

Many different models were considered and tested during the project. The main models that were evaluated include Random Forest, Extra Trees, LightGBM, XGBoost, CatBoost, and neural network models implemented using AutoGluon. Each model has its strengths and weaknesses, and their performance can vary depending on the specific characteristics of the dataset and the problem at hand. [Explain more.](#)

Random Forest and Extra Trees Random Forest and Extra Trees are ensemble learning methods that combine multiple decision trees to improve predictive performance and reduce overfitting.

CatBoost CatBoost is a gradient boosting algorithm that was found to perform well with particular feature sets and hyperparameter configurations [32]. Random Forest and Extra Trees were generally favoured throughout the project, but CatBoost provided competitive results in some scenarios. Since

XGBoost and LightGBM Need to check this out - haven't really given them a real chance with hyperparameter tuning and such yet, as they are suspected to not be ideal for contiguous time-series data.

4.5.3 Adjusting Classification Thresholds

REDUNDANT AS WE ONLY CONSIDER 3-CLASS CLASSIFICATION. A critical aspect of optimizing classification models involves adjusting the decision thresholds to balance precision and recall effectively. It is often useful to find the threshold that maximizes the F1-score, which is the harmonic mean of precision and recall. This approach ensures that both false positives and false negatives are minimized, leading to a more balanced model performance. This was implemented by evaluating the model's performance across a range of thresholds and selecting the one that yielded the highest F1-score. A typical precision-recall curve is shown in figure 13. During training and model evaluation, this threshold was heavily utilized as a general indicator of the model's performance. This streamlined the evaluation process, making it easier to compare different models and configurations based on a single metric. It is important to note that while maximizing the F1-score provides a balanced approach, this is not necessarily the optimal strategy threshold for all applications. Depending on the specific use case, one might prioritize either precision or recall more heavily.

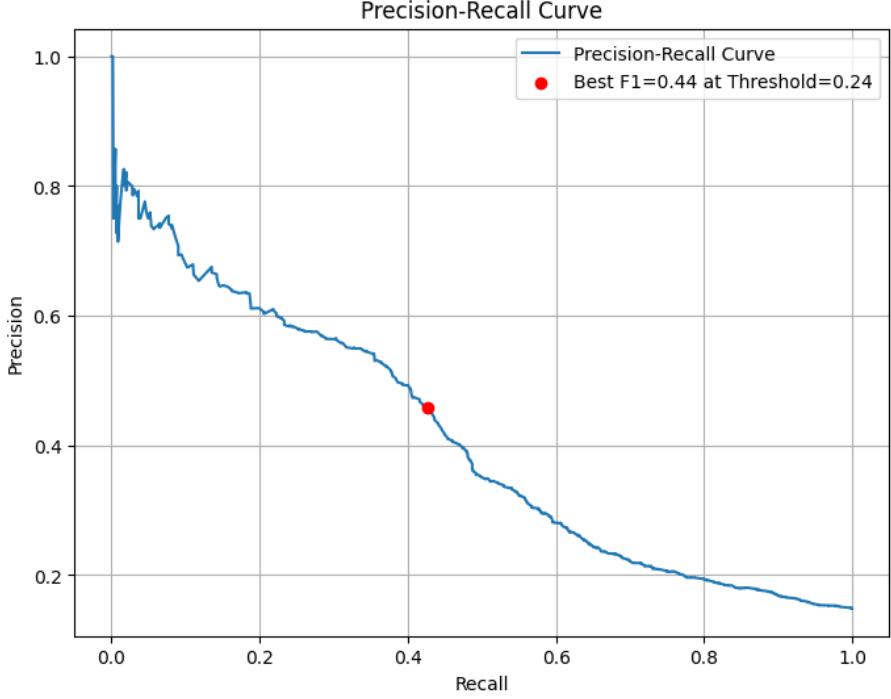


Figure 13: Typical Precision-Recall Curve with F1-score maximization point indicated.

4.5.4 Decision-bias tuning

After extending the prediction task from binary (up vs. none) to ternary (up, down, none), the model performance on the “up” class dropped substantially. This was expected, as discussed in Section 4.3.4, because “up” activations are rare and therefore provide the model with far fewer examples to learn from. As a result, the model became overly conservative and almost always favoured predicting the majority classes.

To address this issue, a simple post-processing method—here referred to as *decision-bias tuning*—was applied. The idea is similar in spirit to adjusting the decision threshold in binary classification, but adapted to the three-class setting. Instead of changing the model itself, we gently shift its decision boundary so that it becomes more willing to predict the “up” class.

The procedure is straightforward. After the model outputs its class probabilities, the predicted probability for the “up” class is multiplied by a bias factor (greater than one). This does not change the underlying model or its learned parameters; it only adjusts how the final class label is chosen. If the model was previously hesitating to predict “up” even when it assigned a moderate probability to it, this adjustment makes such predictions more likely.

The strength of the bias is selected using the validation dataset. Several candidate bias levels are tested, and for each of them we compute the resulting F1-score for the “up” class by keeping the other classes unchanged. The bias factor that gives the best validation performance is then used when making predictions on the test set. This tuning process helps correct the model’s natural tendency to favour the majority classes and provides a controlled way of improving recall and overall F1-score for the under-represented “up” class, without degrading the performance on the remaining classes more than necessary.

Intuition and benefits. This method effectively lowers the decision threshold for predicting the “up” class, making the model more inclined to predict “up” when there is uncertainty. This will improve recall for the “up” class, at the cost of precision. By selecting α based on F1-macro score, we ensure that the trade-off between precision and recall across all classes is optimized for practical performance.

4.5.5 AutoGluon

5 Model Results

During the project, various models were trained on different dataset iterations as new data was added and features were engineered and refined. Models were initially trained on data from 2025 only. This dataset served as the starting point, excluding data from 2024 for simplicity and to establish a baseline. Using a one-year dataset is inherently problematic due to the limited amount of data and the fact that the model can not learn from seasonal patterns. It is also likely that the model overfits to specific events in the year, which may not generalize well to other years. Later, data from 2024 was included to provide more training data and to allow the model to learn from seasonal variations. A two-year dataset is, however, still problematic, as it only captures two instances of seasonal patterns, which is likely to not be representative of previous or coming years. Ideally, the dataset would span multiple years to capture a wider range of seasonal patterns and outliers, but this was not possible due to data availability constraints.

The problem started as a binary classification problem, where the model's task was to predict whether an up-activation would occur in a given hour or not. However, as the project progressed, it was expanded to also include the prediction of down-activations, making it a multi-class classification problem with three classes: up-activation, down-activation, and no activation. The results presented here touch on both the binary and multi-class classification problems, but focuses primarily on the multi-class problem, as it is the most comprehensive and relevant to the real-world application.

5.1 Naive Model

Naive model predicts the same class for $t+4$ as $t-4$. Sadly, with current models - this approach seems to be performing better or similarly to advanced models. 40.63%, 60.35%, 63.59% for up, down, no-activation respectively.

5.2 The 2025 Dataset

Depending on the specific metrics used for evaluation, different models perform best. It was quite early on discovered that Random Forest and Extra Trees models performed best overall, so the results focus on these two models. The models were trained and evaluated based on F1-score, with raw precision and recall also considered after training. The aforementioned transition metric is also used to evaluate the models, as it provides insight into the model's ability to predict activation transitions specifically.

5.2.1 Extra Trees

Figure 14 shows the Precision-Recall curve for an Extra Trees model trained on the 2025 dataset. Three dots are scattered on the curve, representing the precision-recall pairs for three interesting thresholds: the threshold given by the predictor leaderboard, the threshold that maximizes F1-score, and the threshold that gives a recall of 0.5. It is interesting that the leaderboard threshold gives a lower F1-score than the maximum F1-score threshold after training, as one would expect the leaderboard threshold to be optimal. Nevertheless, all three thresholds provide solid performance, with the maximum F1-score threshold achieving the best balance between precision and recall. The threshold that gives a recall of 0.5 sacrifices some precision to achieve higher recall, while simultaneously improving the transition metric.

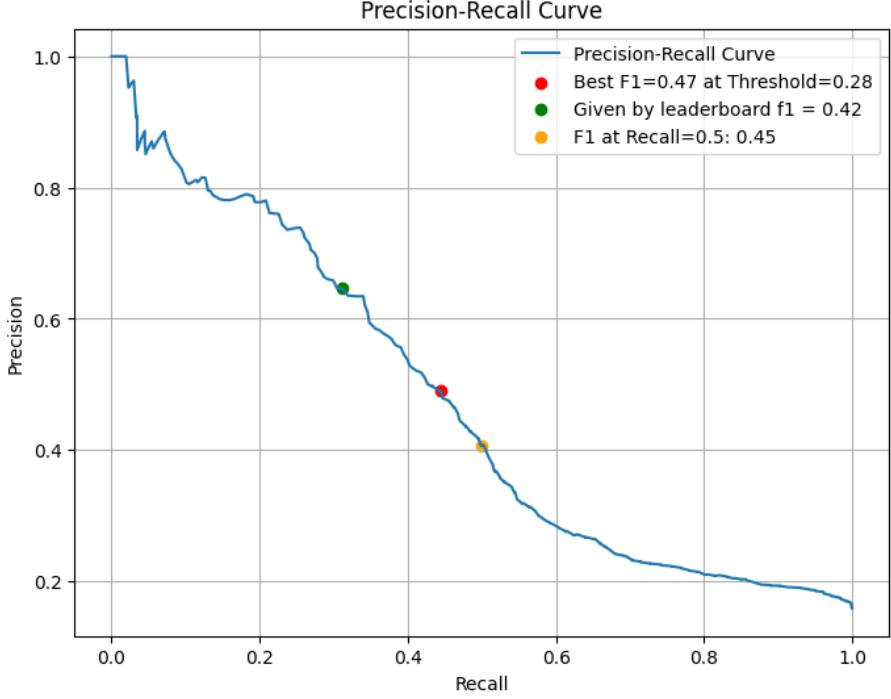


Figure 14: Precision-Recall Curve for Extra Trees model trained on 2025 dataset with highest F1-score.

Threshold	Precision	Recall	F1-score	Transition Metric
Leaderboard	0.65	0.31	0.42	9.49%
Max F1-score	0.49	0.45	0.47	19.76%
Recall = 0.5	0.41	0.50	0.45	28.06%

Table 3: Performance metrics for Extra Trees model on 2025 dataset at different thresholds.

Table 3 summarizes the performance metrics for the Extra Trees model on the 2025 dataset at the three different thresholds. The transition metric is highly correlated with recall, as expected, since higher recall means more activation events are correctly identified, leading to better transition detection. The recall-focused threshold achieves a transition success rate of 28.06%, significantly higher than the threshold proposed by the leaderboard, and 10% higher than the maximum F1-score threshold. With a precision of 0.41 at this threshold, the model is relatively precise, not overwhelmingly predicting activations, which is crucial.

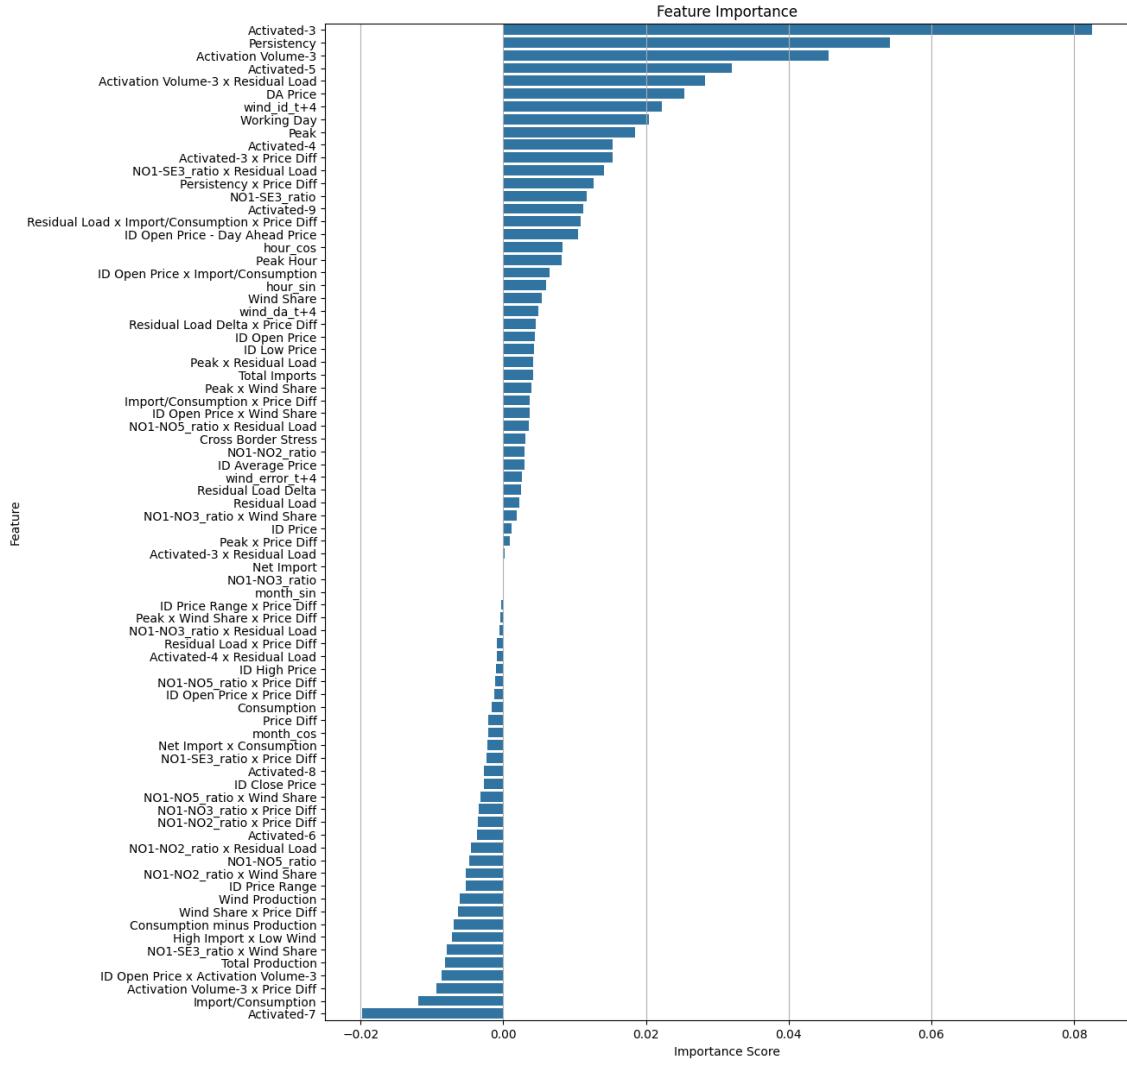


Figure 15: Feature importance for Extra Trees model trained on 2025 dataset with highest F1-score.

Figure 15 shows the feature importance for the Extra Trees model. An important concept to note is that feature importance is independent of the chosen threshold. The feature importance indicates which features the model contribute the most to the underlying probability estimates. Threshold adjustments merely shift the decision boundary without altering the relative importance of the features. The top three features are in this particular model all lag features, specifically persistency and singular activation lag features. This suggests that the model relies heavily on historical activation patterns to make its predictions. Although one of the goals was to reduce the reliance on lag features, it is still important to catch on to activation trends. Other important features include day-ahead price, intraday wind forecasts, time-related features such as peak hour and working day indicators, and various interaction features. All these features likely contribute to capturing the complex dynamics influencing activation events.

5.3 The 2024-2025 Dataset

Some extra preprocessing is necessary when including data from 2024, as data often come in yearly batches. Separate CSV (Comma Separated Values) files for 2024 and 2025 are, therefore, merged into a single dataset before further handling. This dataset was used for most of the training and evaluation process, as it provides more data for the models to learn from, potentially leading to better generalization and performance. This comes at the cost of slightly inconsistent data.

mFRR activation data, for instance, transitioned from hourly to 15-minute resolution in mid-2024 (**specify date perhaps**). This introduces inconsistency and noise into the dataset, which could affect model performance. However, the benefits of having a larger dataset likely outweigh the this drawback.

5.3.1 CatBoost Models

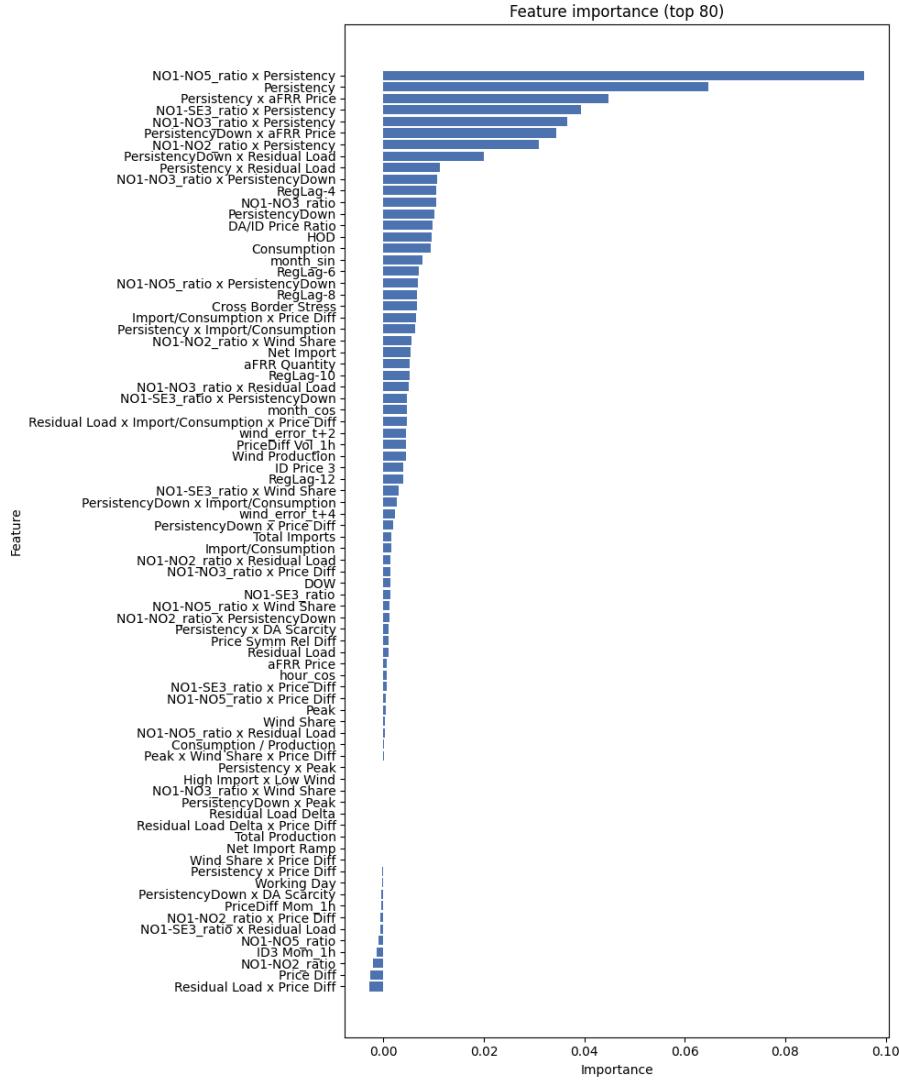


Figure 16: Feature importance for CatBoost model trained on 2024-2025 dataset after hyperparameter tuning.

I probably want a kind of sweep over performances on different hyperparameter combinations.

Metrics table

	val_f1_macro	val_accuracy	test_f1_macro	test_accuracy
0	0.633974	0.714051	0.59821	0.649165

Val classification report

	precision	recall	f1-score	support
down	0.73	0.71	0.72	4957
none	0.74	0.76	0.75	6274
up	0.44	0.42	0.43	918
accuracy			0.71	12149
macro avg	0.64	0.63	0.63	12149
weighted avg	0.71	0.71	0.71	12149

Test classification report

	precision	recall	f1-score	support
down	0.72	0.62	0.67	5112
none	0.63	0.73	0.68	5603
up	0.49	0.41	0.45	1436
accuracy			0.65	12151
macro avg	0.61	0.59	0.60	12151
weighted avg	0.65	0.65	0.65	12151

Figure 17: Performance metrics for CatBoost model on 2024-2025 dataset at different thresholds.

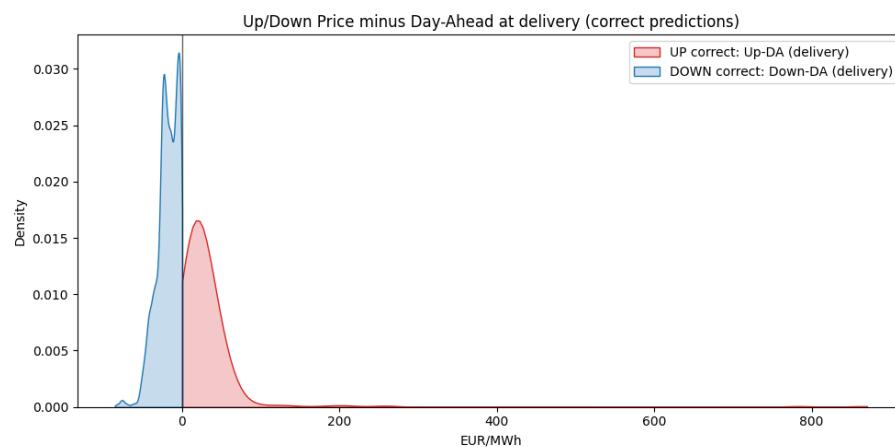


Figure 18: Up/down price minus day-ahead price distribution for CatBoost model trained on 2024-2025 dataset. quick_multiclass_cat_hpo

Figure 19 shows the performance metrics for a CatBoost model trained on data from March 4th, 2025. The metrics are evaluated at specific dataset subsets determined by how confident the model is in its predictions. For instance, at a confidence threshold of 0.6, only predictions where the model's predicted probability for the chosen class is at least 0.6 are considered. This approach allows for an analysis of how the model's performance varies with its confidence level. As the confidence threshold increases, accuracy steadily improves. This is expected, as higher confidence predictions should generally be more reliable. F1-macro score dips, however, at higher thresholds, especially beyond 0.7. The most likely reason for this is that the model rarely is very confident in predicting the less frequent classes (up- and down-activations). Most of these predictions are no-activation predictions, resulting in good accuracy (0.778 in this case). Assume that all of these predictions are no-activation predictions (**can probably check this quickly**). Then, 22.2% of the predictions are false negatives for the up- and down-activation classes, leading to low recall (0 in this case) and thus low F1-score for these classes. The overall F1-macro score, being the average of the F1-scores for all classes, consequently drops as well.

threshold	coverage	acc	f1_macro
0	0.4	0.683801	0.648802
1	0.5	0.711508	0.655303
2	0.6	0.742613	0.657654
3	0.7	0.751748	0.613787
4	0.8	0.778443	0.291807
5	0.9	NaN	NaN

Figure 19: Performance metrics for CatBoost model on 2025 March 4th dataset at different confidence thresholds.

5.4 Correlation

I do not know where this will end up in the structure, but talking about correlation between features is key. Analysis has been done on correlation using different methods. Explaining corelation.

Bibliography

- [1] X. Cai, N. Zhang, E. Du, Z. An, N. Wei and C. Kang, ‘Low Inertia Power System Planning Considering Frequency Quality Under High Penetration of Renewable Energy’, *IEEE Transactions on Power Systems*, vol. PP, pp. 1–12, Jan. 2023. DOI: 10.1109/TPWRS.2023.3302515
- [2] ENTSO-e, *Nordic Balancing Philosophy ENTSOE*, 2024. Accessed: 13th Nov. 2025.
- [3] *Confirmation of mFRR EAM go live March 4th 2025*, <https://www.statnett.no/en/for-stakeholders-in-the-power-industry/news-for-the-power-industry/confirmation-of-mfrr-eam-go-live-march-4th-2025/>, Oct. 2025. Accessed: 22nd Nov. 2025.
- [4] *The Nordic power market introduces 15-minute balancing*, <https://www.value.com/news/nordic-power-market-introduces-15min-balancing>. Accessed: 9th Dec. 2025.
- [5] *Energy Transition Outlook Norway 2024*, <https://www.norskindustri.no/siteassets/dokumenter/rapporter-og-brosjyrer/energy-transition-norway/energy-transition-norway-2024.pdf>. Accessed: 3rd Dec. 2025.
- [6] *Raske frekvensreserver - FFR*, <https://www.statnett.no/for-aktorer-i-kraftbransjen/systemansvaret/kraftmarkedet/reservemarkeder/ffr/>, Nov. 2025. Accessed: 13th Nov. 2025.
- [7] Statnett, *Vilkår for mFRR aktiveringsmarked*, Jan. 2024.
- [8] L. Zhao, ‘Event Prediction in the Big Data Era: A Systematic Survey’, *ACM Computing Surveys*, vol. 54, no. 5, pp. 1–37, Jun. 2022, ISSN: 0360-0300, 1557-7341. DOI: 10.1145/3450287 Accessed: 8th Dec. 2025.
- [9] M. Håberg and G. Doorman, ‘A stochastic mixed integer linear programming formulation for the balancing energy activation problem under uncertainty’, in *2017 IEEE Manchester PowerTech*, Jun. 2017, pp. 1–6. DOI: 10.1109/PTC.2017.7980980 Accessed: 9th Dec. 2025.
- [10] A. Khodadadi, H. Nordström, R. Eriksson and L. Söder, ‘Investigating Reserve Dimensioning Approaches for Multi-Area Reserve Capacity Markets: A Nordic Case Study’, Accessed: 9th Dec. 2025.
- [11] A. Khodadadi and L. Söder, ‘On the Optimal Coordinated Hydropower Bidding Strategy in Day-Ahead Energy and Manual Frequency Restoration Reserve Markets’, in *2021 IEEE Madrid PowerTech*, Jun. 2021, pp. 1–6. DOI: 10.1109/PowerTech46648.2021.9494813 Accessed: 9th Dec. 2025.
- [12] G. Klæboe, J. Braathen, A. L. Eriksrud and S.-E. Fleten, ‘Day-ahead market bidding taking the balancing power market into account’, *TOP*, vol. 30, no. 3, pp. 683–703, Oct. 2022, ISSN: 1134-5764, 1863-8279. DOI: 10.1007/s11750-022-00645-1 Accessed: 1st Dec. 2025.
- [13] ‘Balancing market outlook 2030’,
- [14] *Transition to 15-minute Market Time Unit (MTU)*, <https://www.nordpoolgroup.com/en/trading/transition-to-15-minute-market-time-unit-mtu/>. Accessed: 2nd Dec. 2025.
- [15] V. V. Kallset and H. Farahmand, ‘Improving Balancing Activation Through Continuous-Time Optimization and Increased Market Time-Resolution’, in *2025 21st International Conference on the European Energy Market (EEM)*, May 2025, pp. 1–6. DOI: 10.1109/EEM64765.2025.11050190 Accessed: 2nd Dec. 2025.
- [16] C. Singh, S. Sreekumar and T. Malakar, ‘A novel dynamic imbalance volume forecasting model for balancing market optimization’, *Electrical Engineering*, vol. 107, no. 12, pp. 15 375–15 392, Dec. 2025, ISSN: 1432-0487. DOI: 10.1007/s00202-025-03331-0 Accessed: 2nd Dec. 2025.
- [17] K. Plakas, N. Andriopoulos, D. Papadaskalopoulos, A. Birbas, E. Housos and I. Moraitis, ‘Prediction of Imbalance Prices Through Gradient Boosting Algorithms: An Application to the Greek Balancing Market’, *IEEE Access*, vol. 13, pp. 103 968–103 981, 2025, ISSN: 2169-3536. DOI: 10.1109/ACCESS.2025.3580274 Accessed: 3rd Dec. 2025.
- [18] D. Azarang and C. Edling, ‘Machine Learning-Based Prediction and Key Drivers of mFRR Activations’,

-
- [19] E. R. A. Overmaat, ‘Balancing Contributions in the Nordic Electricity System’,
 - [20] T. Svedlindh and K. Yngvesson, *Price Formation and Forecasting Models in the Electricity Market : An Analysis of the Intraday and mFRR Markets*. 2025. Accessed: 29th Nov. 2025.
 - [21] R. C. Porras, ‘Short-Term Forecasting of mFRR Activation Direction and Imbalance Price using XGBoost’,
 - [22] *Power Market Data*, <https://www.nordpoolgroup.com/en/services/power-market-data-services/>. Accessed: 22nd Nov. 2025.
 - [23] *Static Content — Nordic Unavailability Collection System*, https://www.nucs.net/content/static_content/Static%20content/data%20repository/DataRepositoryGuide.html. Accessed: 22nd Nov. 2025.
 - [24] *Data & Standardisation*, <https://www.entsoe.eu/data/>. Accessed: 22nd Nov. 2025.
 - [25] *Transition to 15-minute Market Time Unit (MTU)*, <https://www.nordpoolgroup.com/en/trading/transition-to-15-minute-market-time-unit-mtu/>. Accessed: 8th Dec. 2025.
 - [26] *Class Imbalance Problem - an overview — ScienceDirect Topics*, <https://www.sciencedirect.com/topics/computer-science/class-imbalance-problem>. Accessed: 26th Nov. 2025.
 - [27] H. Liu and H. Motoda, *Feature Selection for Knowledge Discovery and Data Mining*. Springer Science & Business Media, Dec. 2012, ISBN: 978-1-4615-5689-3.
 - [28] J. Brownlee, *Basic Feature Engineering With Time Series Data in Python*, Dec. 2016. Accessed: 22nd Nov. 2025.
 - [29] H. Pelletier, *Cyclical Encoding: An Alternative to One-Hot Encoding for Time Series Features*, May 2024. Accessed: 22nd Nov. 2025.
 - [30] *Classification: Accuracy, recall, precision, and related metrics — Machine Learning*, <https://developers.google.com/machine-learning/crash-course/classification/accuracy-precision-recall>. Accessed: 22nd Nov. 2025.
 - [31] M. Grandini, E. Bagli and G. Visani, *Metrics for Multi-Class Classification: An Overview*, Aug. 2020. DOI: 10.48550/arXiv.2008.05756 arXiv: 2008.05756 [stat]. Accessed: 9th Dec. 2025.
 - [32] A. V. Dorogush, V. Ershov and A. Gulin, *CatBoost: Gradient boosting with categorical features support*, Oct. 2018. DOI: 10.48550/arXiv.1810.11363 arXiv: 1810.11363 [cs]. Accessed: 22nd Nov. 2025.